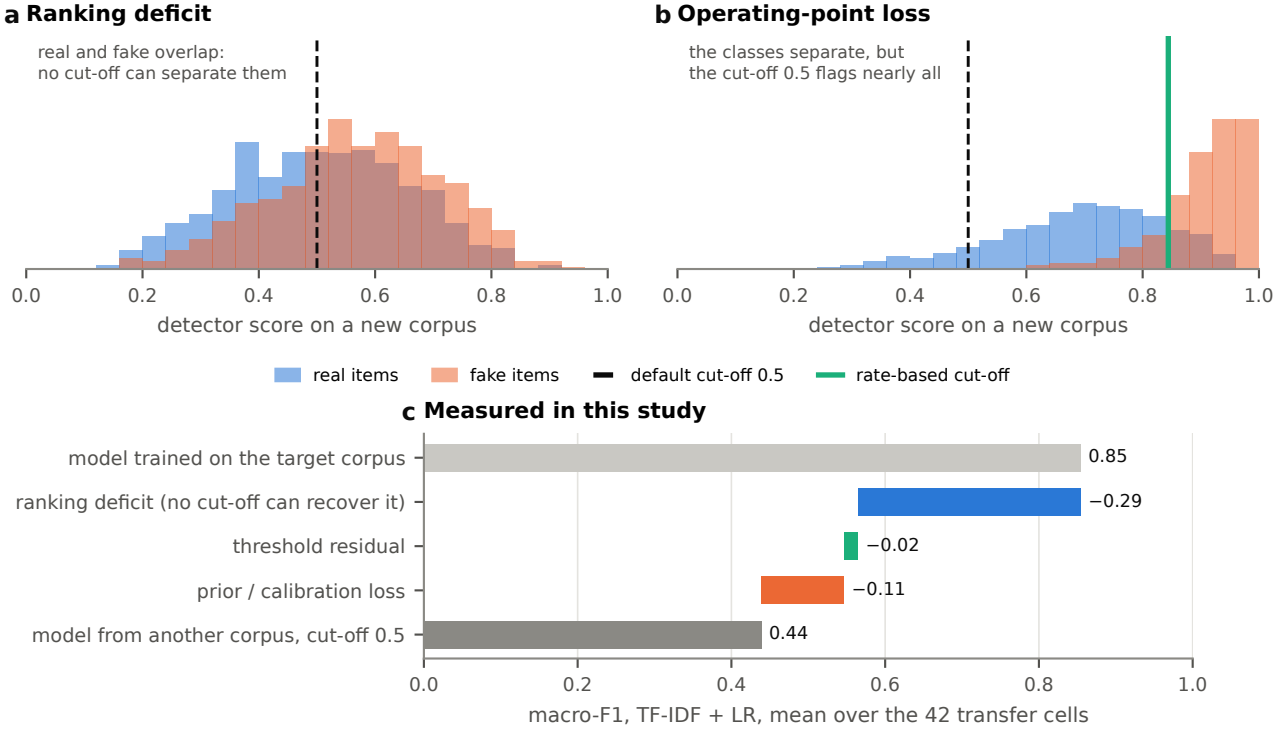
a Ranking deficit
real and fake overlap:
no cut-off can separate them
0.0
0.2
0.4
0.6
0.8
1.0
detector score on a new corpus
b Operating-point loss
the classes separate, but
the cut-off 0.5 flags nearly all
0.0
0.2
0.4
0.6
0.8
1.0
detector score on a new corpus
real items
fake items
default cut-off 0.5
rate-based cut-off
c Measured in this study
model trained on the target corpus
0.85
ranking deficit (no cut-off can recover it)
−0.29
threshold residual
−0.02
prior / calibration loss
−0.11
model from another corpus, cut-off 0.5
0.44
0.0
0.2
0.4
0.6
0.8
1.0
macro-F1, TF-IDF + LR, mean over the 42 transfer cells