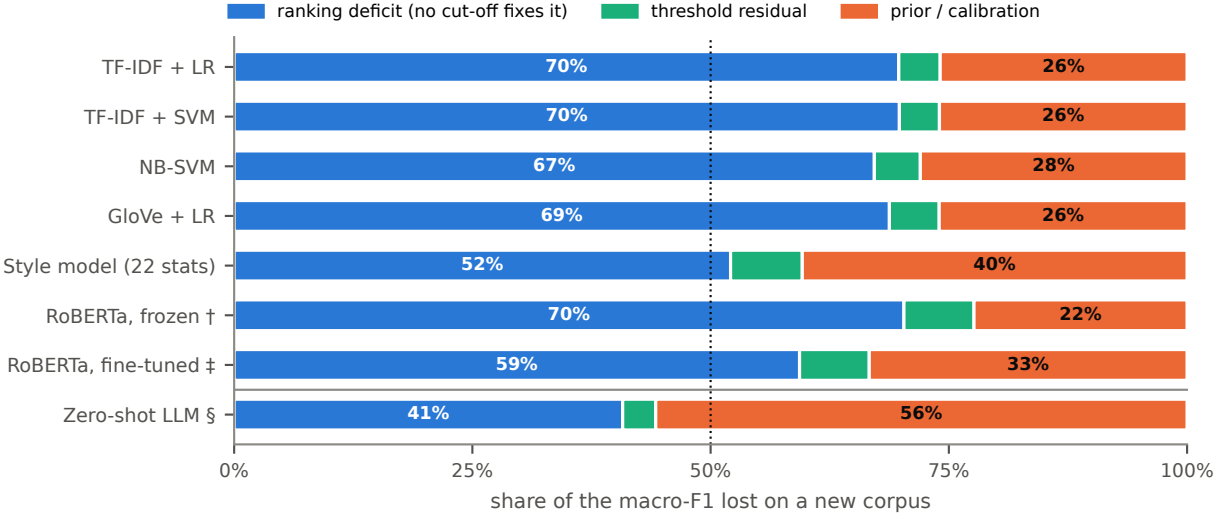
ranking deficit (no cut-off fixes it)
threshold residual
prior / calibration
TF-IDF + LR
TF-IDF + SVM
NB-SVM
GloVe + LR
Style model (22 stats)
RoBERTa, frozen †
RoBERTa, fine-tuned ‡
Zero-shot LLM §
70%
26%
70%
26%
67%
28%
69%
26%
52%
40%
70%
22%
59%
33%
41%
56%
0%
25%
50%
75%
100%
share of the macro-F1 lost on a new corpus